

ECON 573 - Final Proejct

```
library(readr)
library(stringr)
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 4.2.2
```

```
library(lubridate)
```

```
##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
```

```
library(boot)
library(caret)
```

```
## Loading required package: lattice
```

```
##
## Attaching package: 'lattice'
```

```
## The following object is masked from 'package:boot':
##
##   melanoma
```

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-4
```

```
library(rpart)
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 4.2.2
```

```
library(ipred)
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##     margin
```

```
library(gbm)
```

```
## Loaded gbm 2.1.8.1
```

```
library(e1071)
```

```
library(MASS)
```

```
library(Rfast)
```

```
## Warning: package 'Rfast' was built under R version 4.2.2
```

```
## Loading required package: Rcpp
```

```
## Loading required package: RcppZiggurat
```

```
## Warning: package 'RcppZiggurat' was built under R version 4.2.2
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:Rfast':
```

```
##
```

```
##     nth
```

```
## The following object is masked from 'package:MASS':
```

```
##
```

```
##     select
```

```
## The following object is masked from 'package:randomForest':
```

```
##
```

```
##     combine
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
KS_Raw = read.csv("Data/kickstarter_data_full.csv", stringsAsFactors = TRUE)
```

```
# Conversion to USD for FX, Adjusting to Feb 2017 CPI
monthly_cpi = read.table("http://research.stlouisfed.org/fred2/data/CPIAUCSL.txt",
                          skip = 54, header = TRUE)
monthly_cpi$DATE = as.Date(monthly_cpi$DATE)
```

```
KS_Adjusted = KS_Raw %>%
  mutate(goal_usd = goal*static_usd_rate,
         pledged_usd = pledged*static_usd_rate,
         launched_at = as.Date(launched_at),
         yr_mth_launch = floor_date(launched_at, "month")) %>%
  left_join(monthly_cpi, by = c("yr_mth_launch" = "DATE")) %>%
  rename(CPI = VALUE) %>%
  mutate(goal_adj = goal_usd * max(CPI) / CPI,
         pledged_adj = pledged * max(CPI) / CPI)
```

```
# Cleaning
KS_Clean = KS_Adjusted %>%
  select(state, goal_adj, country, category, name_len, name_len_clean, blurb_len,
         blurb_len_clean, deadline_weekday, created_at_weekday, launched_at_weekday,
         deadline_month, deadline_yr, created_at_month, created_at_yr,
         launched_at_month, launched_at_yr, create_to_launch_days,
         launch_to_deadline_days, backers_count, pledged_adj) %>%
  filter(state == "successful" | state == "failed") %>%
  mutate(success = ifelse(state == "successful", 1, 0),
         country = relevel(country, ref = "US"),
         deadline_month = as.factor(deadline_month),
         deadline_yr = as.factor(deadline_yr),
         created_at_month = as.factor(created_at_month),
         created_at_yr = as.factor(created_at_yr),
         launched_at_month = as.factor(launched_at_month),
         launched_at_yr = as.factor(launched_at_yr),
         avg_pledge = ifelse(backers_count == 0, 0, pledged_adj/backers_count)) %>%
  filter(category != "Comedy", #1
         country != "LU", #2
         ) %>% # Removed for CV issues
  select(success, everything(), -state, -backers_count, -pledged_adj)

KS_Clean$category = as.character(KS_Clean$category)
KS_Clean$category[KS_Clean$category == ""] = "Other"
KS_Clean$category = relevel(as.factor(KS_Clean$category), ref = "Other")
KS_Clean$country = droplevels(KS_Clean$country)

write.csv(KS_Clean, "KS_Clean.csv")
```

```
# Sampling of Test / Train
seed = 5
set.seed(seed)

Train_Ind = sample(1:nrow(KS_Clean),
                  round(.80*nrow(KS_Clean))) # 80:20 Train/Test Split
KS_Train = KS_Clean[Train_Ind,]
KS_Test = KS_Clean[-Train_Ind,]
```

```
# Logistic Regression
Logistic_Model = glm(success ~ ., family = binomial, data = KS_Train)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(Logistic_Model)
```

```
##
## Call:
## glm(formula = success ~ ., family = binomial, data = KS_Train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -5.9895  -0.8353  -0.4062   0.9485   6.4823
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    8.866e-01  5.140e-01   1.725 0.084535 .
## goal_adj      -9.925e-06  6.068e-07 -16.357 < 2e-16 ***
## countryAT     -8.993e-01  4.580e-01  -1.963 0.049593 *
## countryAU     -3.809e-01  1.336e-01  -2.852 0.004346 **
## countryBE     -1.288e+00  6.502e-01  -1.981 0.047613 *
## countryCA     -2.995e-01  9.832e-02  -3.046 0.002318 **
## countryCH     -1.003e-01  3.329e-01  -0.301 0.763126
## countryDE      1.911e-02  1.614e-01   0.118 0.905725
## countryDK     -1.350e+00  3.985e-01  -3.387 0.000706 ***
## countryES     -9.735e-01  2.984e-01  -3.262 0.001106 **
## countryFR     -1.819e-01  1.870e-01  -0.973 0.330669
## countryGB      1.715e-01  6.348e-02   2.702 0.006891 **
## countryHK     -1.305e+00  6.562e-01  -1.988 0.046783 *
## countryIE      2.381e-01  3.262e-01   0.730 0.465586
## countryIT     -1.236e+00  2.772e-01  -4.460 8.19e-06 ***
## countryMX     -3.391e+00  7.045e-01  -4.814 1.48e-06 ***
## countryNL     -1.060e-01  1.770e-01  -0.599 0.549462
## countryNO     -2.049e+00  5.887e-01  -3.480 0.000502 ***
## countryNZ     -1.568e-01  2.975e-01  -0.527 0.598043
## countrySE     -1.564e+00  3.853e-01  -4.059 4.92e-05 ***
## countrySG     -2.053e-01  6.156e-01  -0.333 0.738786
## categoryAcademic -1.506e+01  3.646e+02  -0.041 0.967060
## categoryApps   -5.787e-01  1.023e-01  -5.659 1.52e-08 ***
## categoryBlues   1.541e+01  3.795e+02   0.041 0.967610
## categoryExperimental 8.137e-01  1.537e-01   5.295 1.19e-07 ***
## categoryFestivals 8.289e-01  1.339e-01   6.192 5.93e-10 ***
```

## categoryFlight	-1.056e+00	1.774e-01	-5.956	2.59e-09	***
## categoryGadgets	-3.270e-01	8.445e-02	-3.872	0.000108	***
## categoryHardware	-4.481e-01	7.892e-02	-5.678	1.36e-08	***
## categoryImmersive	5.904e-01	1.628e-01	3.626	0.000288	***
## categoryMakerspaces	-1.673e-01	1.957e-01	-0.855	0.392609	
## categoryMusical	4.495e-01	1.087e-01	4.136	3.53e-05	***
## categoryPlaces	-1.523e+01	1.582e+02	-0.096	0.923329	
## categoryPlays	6.713e-01	9.661e-02	6.948	3.71e-12	***
## categoryRobots	4.209e-02	1.364e-01	0.309	0.757602	
## categoryShorts	1.514e+01	2.595e+02	0.058	0.953473	
## categorySoftware	-1.588e+00	9.376e-02	-16.938	< 2e-16	***
## categorySound	8.480e-02	1.311e-01	0.647	0.517634	
## categorySpaces	4.284e-01	1.997e-01	2.145	0.031916	*
## categoryThrillers	-1.475e+01	3.405e+02	-0.043	0.965443	
## categoryWearables	-1.212e-01	1.082e-01	-1.120	0.262661	
## categoryWeb	-1.915e+00	1.006e-01	-19.032	< 2e-16	***
## categoryWebseries	-1.562e+01	3.545e+02	-0.044	0.964857	
## name_len	7.917e-02	2.141e-02	3.698	0.000217	***
## name_len_clean	4.518e-02	2.522e-02	1.791	0.073219	.
## blurb_len	-4.067e-02	7.491e-03	-5.429	5.66e-08	***
## blurb_len_clean	3.959e-02	1.060e-02	3.734	0.000189	***
## deadline_weekdayMonday	1.380e-01	8.472e-02	1.628	0.103462	
## deadline_weekdaySaturday	-2.104e-02	8.055e-02	-0.261	0.793975	
## deadline_weekdaySunday	-7.623e-02	7.952e-02	-0.959	0.337787	
## deadline_weekdayThursday	-3.895e-02	7.780e-02	-0.501	0.616596	
## deadline_weekdayTuesday	2.062e-01	8.786e-02	2.346	0.018953	*
## deadline_weekdayWednesday	4.892e-02	7.891e-02	0.620	0.535305	
## created_at_weekdayMonday	1.700e-02	7.581e-02	0.224	0.822545	
## created_at_weekdaySaturday	-1.368e-01	8.756e-02	-1.562	0.118171	
## created_at_weekdaySunday	-1.624e-01	8.537e-02	-1.902	0.057184	.
## created_at_weekdayThursday	-4.671e-02	7.774e-02	-0.601	0.547972	
## created_at_weekdayTuesday	9.961e-02	7.531e-02	1.323	0.185906	
## created_at_weekdayWednesday	-1.007e-02	7.699e-02	-0.131	0.895882	
## launched_at_weekdayMonday	2.043e-01	8.328e-02	2.453	0.014161	*
## launched_at_weekdaySaturday	-1.879e-02	1.121e-01	-0.168	0.866906	
## launched_at_weekdaySunday	1.852e-01	1.106e-01	1.675	0.093981	.
## launched_at_weekdayThursday	1.507e-01	8.590e-02	1.754	0.079356	.
## launched_at_weekdayTuesday	4.849e-01	8.106e-02	5.982	2.21e-09	***
## launched_at_weekdayWednesday	2.668e-01	8.181e-02	3.261	0.001111	**
## deadline_month2	3.026e-01	1.898e-01	1.594	0.110863	
## deadline_month3	3.762e-01	2.668e-01	1.410	0.158506	
## deadline_month4	3.084e-01	3.342e-01	0.923	0.356077	
## deadline_month5	4.891e-01	3.965e-01	1.234	0.217351	
## deadline_month6	5.965e-01	4.533e-01	1.316	0.188176	
## deadline_month7	5.531e-01	5.088e-01	1.087	0.277046	
## deadline_month8	6.477e-01	5.651e-01	1.146	0.251679	
## deadline_month9	4.956e-01	6.221e-01	0.797	0.425623	
## deadline_month10	5.878e-01	6.812e-01	0.863	0.388239	
## deadline_month11	7.277e-01	7.379e-01	0.986	0.324053	
## deadline_month12	6.904e-01	7.954e-01	0.868	0.385405	
## deadline_yr2010	6.919e-01	1.185e+00	0.584	0.559165	
## deadline_yr2011	1.057e+00	1.979e+00	0.534	0.593478	
## deadline_yr2012	1.721e+00	2.785e+00	0.618	0.536677	
## deadline_yr2013	2.226e+00	3.601e+00	0.618	0.536421	

```

## deadline_yr2014      3.092e+00  4.415e+00   0.700 0.483726
## deadline_yr2015      3.633e+00  5.237e+00   0.694 0.487802
## deadline_yr2016      4.562e+00  6.069e+00   0.752 0.452276
## deadline_yr2017      5.126e+00  6.903e+00   0.743 0.457774
## created_at_month2    -3.678e-02  1.307e-01  -0.282 0.778313
## created_at_month3      1.384e-01  1.658e-01   0.835 0.403623
## created_at_month4      9.205e-02  2.114e-01   0.436 0.663176
## created_at_month5      1.179e-01  2.595e-01   0.454 0.649474
## created_at_month6      1.269e-01  3.090e-01   0.411 0.681242
## created_at_month7     -1.244e-01  3.639e-01  -0.342 0.732467
## created_at_month8     -1.641e-01  4.175e-01  -0.393 0.694249
## created_at_month9     -1.336e-01  4.724e-01  -0.283 0.777267
## created_at_month10    -3.565e-01  5.270e-01  -0.676 0.498753
## created_at_month11    -4.455e-01  5.796e-01  -0.769 0.442072
## created_at_month12    -4.597e-01  6.369e-01  -0.722 0.470414
## created_at_yr2010     -1.657e+01  1.455e+03  -0.011 0.990916
## created_at_yr2011     -1.839e+01  1.455e+03  -0.013 0.989916
## created_at_yr2012     -1.920e+01  1.455e+03  -0.013 0.989476
## created_at_yr2013     -1.971e+01  1.455e+03  -0.014 0.989193
## created_at_yr2014     -2.034e+01  1.455e+03  -0.014 0.988848
## created_at_yr2015     -2.079e+01  1.455e+03  -0.014 0.988603
## created_at_yr2016     -2.126e+01  1.455e+03  -0.015 0.988345
## created_at_yr2017     -2.258e+01  1.455e+03  -0.016 0.987623
## launched_at_month2      7.581e-03  1.879e-01   0.040 0.967818
## launched_at_month3     -3.654e-02  2.631e-01  -0.139 0.889529
## launched_at_month4     -2.474e-01  3.352e-01  -0.738 0.460611
## launched_at_month5     -2.572e-01  3.939e-01  -0.653 0.513761
## launched_at_month6     -3.217e-01  4.556e-01  -0.706 0.480190
## launched_at_month7     -3.697e-01  5.144e-01  -0.719 0.472330
## launched_at_month8     -1.298e-01  5.742e-01  -0.226 0.821135
## launched_at_month9     -1.861e-01  6.357e-01  -0.293 0.769684
## launched_at_month10    -1.535e-01  6.960e-01  -0.221 0.825397
## launched_at_month11      2.103e-02  7.547e-01   0.028 0.977771
## launched_at_month12    -1.981e-01  8.100e-01  -0.245 0.806797
## launched_at_yr2010      1.497e+01  1.455e+03   0.010 0.991793
## launched_at_yr2011      1.661e+01  1.455e+03   0.011 0.990896
## launched_at_yr2012      1.684e+01  1.455e+03   0.012 0.990768
## launched_at_yr2013      1.674e+01  1.455e+03   0.012 0.990823
## launched_at_yr2014      1.566e+01  1.455e+03   0.011 0.991416
## launched_at_yr2015      1.557e+01  1.455e+03   0.011 0.991462
## launched_at_yr2016      1.521e+01  1.455e+03   0.010 0.991660
## launched_at_yr2017      1.534e+01  1.455e+03   0.011 0.991590
## create_to_launch_days  -1.248e-03  1.887e-03  -0.662 0.508234
## launch_to_deadline_days -1.569e-02  3.005e-03  -5.222 1.77e-07 ***
## avg_pledge            2.223e-03  1.411e-04  15.750 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 18000 on 13944 degrees of freedom
## Residual deviance: 14175 on 13820 degrees of freedom
## AIC: 14425
##

```

```
## Number of Fisher Scoring iterations: 14
```

Logistic Regression Performance


```
Logistic_Pred = ifelse(predict(Logistic_Model, KS_Test, type = "response")>.5, 1, 0)
Logistic_Table = table(Logistic_Pred, KS_Test$success)
Logistic_Table_Prop = prop.table(Logistic_Table)
Logistic_Misc = 1-sum(diag(Logistic_Table_Prop))
Logistic_Misc
```

```
## [1] 0.2507172
```

Lasso Regression

```
x_Train = model.matrix(success~., KS_Train)[,-1]
y_Train = KS_Train$success

set.seed(seed)
CV_Lasso = cv.glmnet(x_Train, y_Train, alpha = 1, family = "binomial", nfolds = 10)
plot(CV_Lasso)
```



Lasso Min

```
Lasso_Model_min = glmnet(x_Train, y_Train, alpha = 1, family = "binomial",
                          lambda = CV_Lasso$lambda.min)

Lasso_Min_Coef = coef(Lasso_Model_min)
```

```

# Lasso 1se
Lasso_Model_1se = glmnet(x_Train, y_Train, alpha = 1, family = "binomial",
                          lambda = CV_Lasso$lambda.1se)

Lasso_1se_Coef = coef(Lasso_Model_1se)

# Comparing Lasso min and Lasso 1se
Lasso_Comparison = data.frame(Vars = Lasso_Min_Coef@Dimnames[[1]],
                              lambda_min_coef = as.numeric(as.character(Lasso_Min_Coef))!=0,
                              lambda_1se_coef = as.numeric(as.character(Lasso_1se_Coef))!=0)

sum(Lasso_Comparison$lambda_min_coef)

## [1] 120

sum(Lasso_Comparison$lambda_1se_coef)

## [1] 99

Lasso_Diff = Lasso_Comparison[which(
  Lasso_Comparison$lambda_min_coef != Lasso_Comparison$lambda_1se_coef),]
Lasso_Diff$Vars

## [1] "countryCH"          "countryDE"
## [3] "countryNZ"            "countrySG"
## [5] "deadline_weekdaySaturday" "deadline_month5"
## [7] "deadline_month8"       "deadline_month12"
## [9] "deadline_yr2012"        "deadline_yr2013"
## [11] "deadline_yr2014"        "deadline_yr2016"
## [13] "created_at_month7"      "created_at_month8"
## [15] "created_at_yr2010"      "created_at_yr2015"
## [17] "created_at_yr2016"      "launched_at_month2"
## [19] "launched_at_month4"     "launched_at_month8"
## [21] "launched_at_month10"    "launched_at_yr2010"
## [23] "launched_at_yr2015"

# Lasso Model Performance - 1se
x_Test = model.matrix(success~., KS_Test)[,-1]

Lasso_Pred = ifelse(predict(Lasso_Model_1se, x_Test, type = "response")>.5, 1, 0)
Lasso_Table = table(Lasso_Pred, KS_Test$success)
Lasso_Table_Prop = prop.table(Lasso_Table)
Lasso_Misc = 1-sum(diag(Lasso_Table_Prop))
Lasso_Misc

## [1] 0.2587493

# Lasso Model Performance - min
x_Test = model.matrix(success~., KS_Test)[,-1]


```



```
Lasso_Pred = ifelse(predict(Lasso_Model_min, x_Test, type = "response")>.5, 1, 0)
Lasso_Table = table(Lasso_Pred, KS_Test$success)
Lasso_Table_Prop = prop.table(Lasso_Table)
Lasso_Misc = 1-sum(diag(Lasso_Table_Prop))
Lasso_Misc # Better
```

```
## [1] 0.2498566
```

```
# Ridge Regression
set.seed(seed)
CV_Ridge = cv.glmnet(x_Train, y_Train, alpha = 0, family = "binomial", nfolds = 10)
plot(CV_Ridge)
```



```
# Ridge Min
Ridge_Model_min = glmnet(x_Train, y_Train, alpha = 1, family = "binomial",
                        lambda = CV_Ridge$lambda.min)
```

```
# Ridge 1se
Ridge_Model_1se = glmnet(x_Train, y_Train, alpha = 1, family = "binomial",
                        lambda = CV_Ridge$lambda.1se)
```

```
# Ridge Model Performance - 1se
Ridge_Pred = ifelse(predict(Ridge_Model_1se, x_Test, type = "response")>.5, 1, 0)
Ridge_Table = table(Ridge_Pred, KS_Test$success)
```

```
Ridge_Table_Prop = prop.table(Ridge_Table)
Ridge_Misc = 1-sum(diag(Ridge_Table_Prop))
Ridge_Misc
```

```
## [1] 0.319564
```

```
# Ridge Model Performance - Min
Ridge_Pred = ifelse(predict(Ridge_Model_min, x_Test, type = "response")>.5, 1, 0)
Ridge_Table = table(Ridge_Pred, KS_Test$success)
Ridge_Table_Prop = prop.table(Ridge_Table)
Ridge_Misc = 1-sum(diag(Ridge_Table_Prop))
Ridge_Misc #Better
```

```
## [1] 0.2908778
```

```
# Decision Tree
set.seed(seed)
Basic_Tree = rpart(success ~ ., data = KS_Train, method = "class",
                    control = rpart.control(cp = 0))


Tree_Min_Error_Ind = which.min(Basic_Tree$cptable[,4])
Tree_Min_Error_Cp = Basic_Tree$cptable[Tree_Min_Error_Ind,1]
```

```
# Basic Tree Performance
Tree_Pred = predict(Basic_Tree, KS_Test, type = "class")

Tree_Table = table(Tree_Pred, KS_Test$success)
Tree_Table_Prop = prop.table(Tree_Table)
Tree_Misc = 1-sum(diag(Tree_Table_Prop))
Tree_Misc
```

```
## [1] 0.2366609
```

```
plot(Basic_Tree$cptable[,1], Basic_Tree$cptable[,4],
     type = "l", col = "black", xlab = "Complexity Parameter", ylab = "CV Error")
abline(v = Basic_Tree$cptable[which.min(Basic_Tree$cptable[,4]),1], col = "red")
```

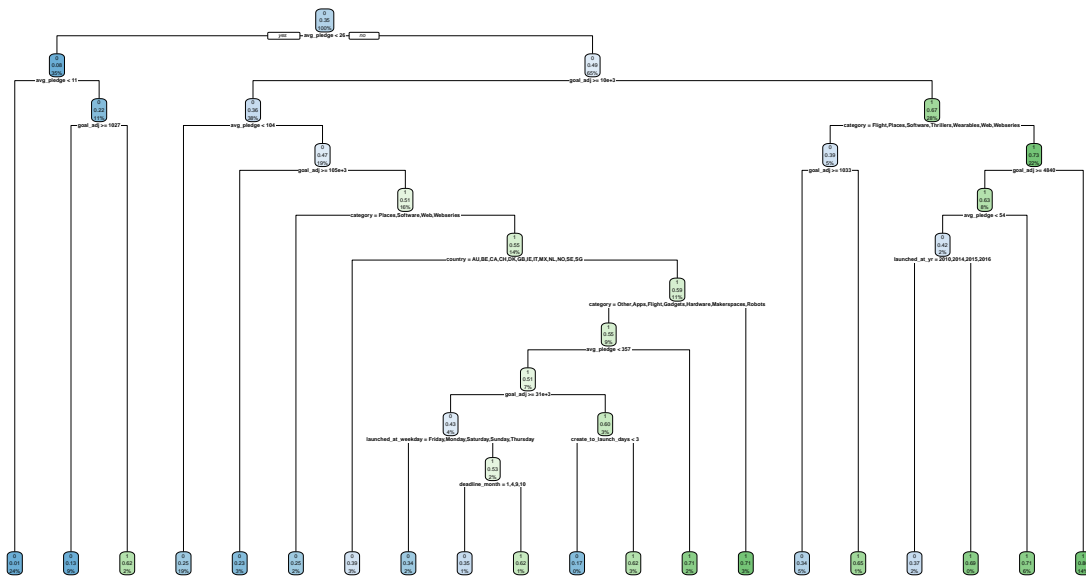


Pruning

```
Pruned_Tree = prune(Basic_Tree, cp = Tree_Min_Error_Cp)
Pruned_Tree$variable.importance
```

```
##          avg_pledge          goal_adj          category
##          1412.4478475          705.9892524          542.8297411
##  create_to_launch_days launch_to_deadline_days          country
##          164.1906871          66.7502067          40.4889536
##          name_len_clean          name_len          deadline_yr
##          40.1689607          39.8291576          17.0368109
##          launched_at_yr          created_at_yr          launched_at_weekday
##          16.6600065          15.6679512          8.4812528
##          deadline_month          launched_at_month          deadline_weekday
##          7.1423730          5.6489952          3.1398681
##          created_at_month          created_at_weekday          blurb_len_clean
##          2.1322749          0.7218087          0.5687519
##          blurb_len
##          0.4420170
```

```
write.csv(Pruned_Tree$variable.importance, "TreeVI.csv")
rpart.plot(Pruned_Tree)
```



Pruned Decision Tree Performance

```
Pruned_Tree_Pred = predict(Pruned_Tree, KS_Test, type = "class")
```

```
Pruned_Tree_Table = table(Pruned_Tree_Pred, KS_Test$success)
```

```
Pruned_Tree_Table_Prop = prop.table(Pruned_Tree_Table)
```

```
Pruned_Tree_Misc = 1-sum(diag(Pruned_Tree_Table_Prop))
```

```
Pruned_Tree_Misc
```

```
## [1] 0.2079748
```

Bagging

```
set.seed(seed)
```

```
Bag_Model = bagging(factor(success) ~ ., data = KS_Train,
```

```
nbagg = 100,
```

```
coob = TRUE,
```

```
control = rpart.control(minsplit = 2, cp = 0))
```

```
Bag_Model
```

```
##
```

```
## Bagging classification trees with 100 bootstrap replications
```

```
##
```

```
## Call: bagging.data.frame(formula = factor(success) ~ ., data = KS_Train,
```

```
## nbagg = 100, coob = TRUE, control = rpart.control(minsplit = 2,
```

```
## cp = 0))
```

```
##
## Out-of-bag estimate of misclassification error: 0.2057

# Bagged Model Variable Importance
VI = varImp(Bag_Model)
VI$var = rownames(VI)
VI_Order = order(VI$Overall, decreasing = TRUE)
VI = VI[VI_Order,]
write.csv(VI, "BaggedVI.csv")

# Bagged Model Performance
Bag_Pred = predict(Bag_Model, KS_Test, type = "class")

Bag_Table = table(Bag_Pred, KS_Test$success)
Bag_Table_Prop = prop.table(Bag_Table)
Bag_Misc = 1-sum(diag(Bag_Table_Prop))
Bag_Misc

## [1] 0.188755

# Random Forest Model
RF_Model = randomForest(factor(success) ~ ., data = KS_Train,
                        ntree = 1000)
RF_Model

##
## Call:
## randomForest(formula = factor(success) ~ ., data = KS_Train,      ntree = 1000)
##           Type of random forest: classification
##           Number of trees: 1000
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 19.83%
## Confusion matrix:
##      0      1 class.error
## 0 7619 1491  0.1636663
## 1 1274 3561  0.2634953

# Random Forest Model Variable Importance
VI_RF = varImp(RF_Model)
VI_RF$var = rownames(VI_RF)
VI_RF_Order = order(VI_RF$Overall, decreasing = TRUE)
VI_RF = VI_RF[VI_RF_Order,]
write.csv(VI_RF, "RFVI.csv")

# Random Forest Performance
RF_Pred = predict(RF_Model, KS_Test, type = "class")

RF_Table = table(RF_Pred, KS_Test$success)
RF_Table_Prop = prop.table(RF_Table)
RF_Misc = 1-sum(diag(RF_Table_Prop))
RF_Misc
```


```
## [1] 0.1864601
```

```
# Boosting
set.seed(seed)
N_Trees = 10000
Boost_Model = gbm(success ~.,
  data = KS_Train,
  distribution = "bernoulli",
  cv.folds = 10,
  shrinkage = .01,
  train.fraction = 0.5,
  n.trees = N_Trees)
Optimal_Trees = which.min(Boost_Model$cv.error)
```

```
# Boosting CV
range = 1:N_Trees
plot(range, Boost_Model$train.error, type = "l", xlab = 'N Trees', ylab = 'CV Error')
lines(range, Boost_Model$cv.error, col = "green")
which.min(Boost_Model$cv.error)
```

```
## [1] 6082
```

```
abline(v = which.min(Boost_Model$cv.error), col = "red")
```



```

# Boost Performance
Boost_Pred = ifelse(predict(Boost_Model,
                           KS_Test, type = "response", n.trees = Optimal_Trees)>.5, 1, 0)

Boost_Table = table(Boost_Pred, KS_Test$success)
Boost_Table_Prop = prop.table(Boost_Table)
Boost_Misc = 1-sum(diag(Boost_Table_Prop))
Boost_Misc

```

```
## [1] 0.1858864
```

```

#SVM
SVM_Linear = svm(factor(success) ~ ., data = KS_Train, kernel = "linear", cost = 5)
SVM_Radial = svm(factor(success) ~ ., data = KS_Train, kernel = "radial", cost = 5)

```

```

# SVM Linear Performance
SVM_Linear_Pred = predict(SVM_Linear, KS_Test, type = "class")

SVM_Linear_Table = table(SVM_Linear_Pred, KS_Test$success)
SVM_Linear_Table_Prop = prop.table(SVM_Linear_Table)
SVM_Linear_Misc = 1-sum(diag(SVM_Linear_Table_Prop))
SVM_Linear_Misc

```

```
## [1] 0.2685026
```

```

# SVM Radial Performance
SVM_Radial_Pred = predict(SVM_Radial, KS_Test, type = "class")

SVM_Radial_Table = table(SVM_Radial_Pred, KS_Test$success)
SVM_Radial_Table_Prop = prop.table(SVM_Radial_Table)
SVM_Radial_Misc = 1-sum(diag(SVM_Radial_Table_Prop))
SVM_Radial_Misc



```

```
## [1] 0.253012
```

```

# LDA
LDA = lda(factor(success) ~.-country -category -deadline_weekday -created_at_weekday -
          launched_at_weekday -deadline_month -deadline_yr -created_at_month -
          created_at_yr -launched_at_yr -launched_at_month, data = KS_Train)
ldahist(data = predict(LDA, KS_Train)$x[,1], g = KS_Train$success)

```



```
# LDA Performance
```

```
LDA_Pred = predict(LDA, KS_Test)$class
```

```
LDA_Table = table(LDA_Pred, KS_Test$success)
```

```
LDA_Table_Prop = prop.table(LDA_Table)
```

```
LDA_Misc = 1-sum(diag(LDA_Table_Prop))
```

```
LDA_Misc
```

```
## [1] 0.3304647
```

```
# QDA
```

```
QDA = qda(factor(success) ~. -country -category -deadline_weekday -created_at_weekday -
  launched_at_weekday -deadline_month -deadline_yr -created_at_month -
  created_at_yr -launched_at_yr -launched_at_month, data = KS_Train)
```

```
# QDA Performance
```

```
QDA_Pred = predict(QDA, KS_Test)$class
```

```
QDA_Table = table(QDA_Pred, KS_Test$success)
```

```
QDA_Table_Prop = prop.table(QDA_Table)
```

```
QDA_Misc = 1-sum(diag(QDA_Table_Prop))
```

```
QDA_Misc
```

```
## [1] 0.6044177
```



```

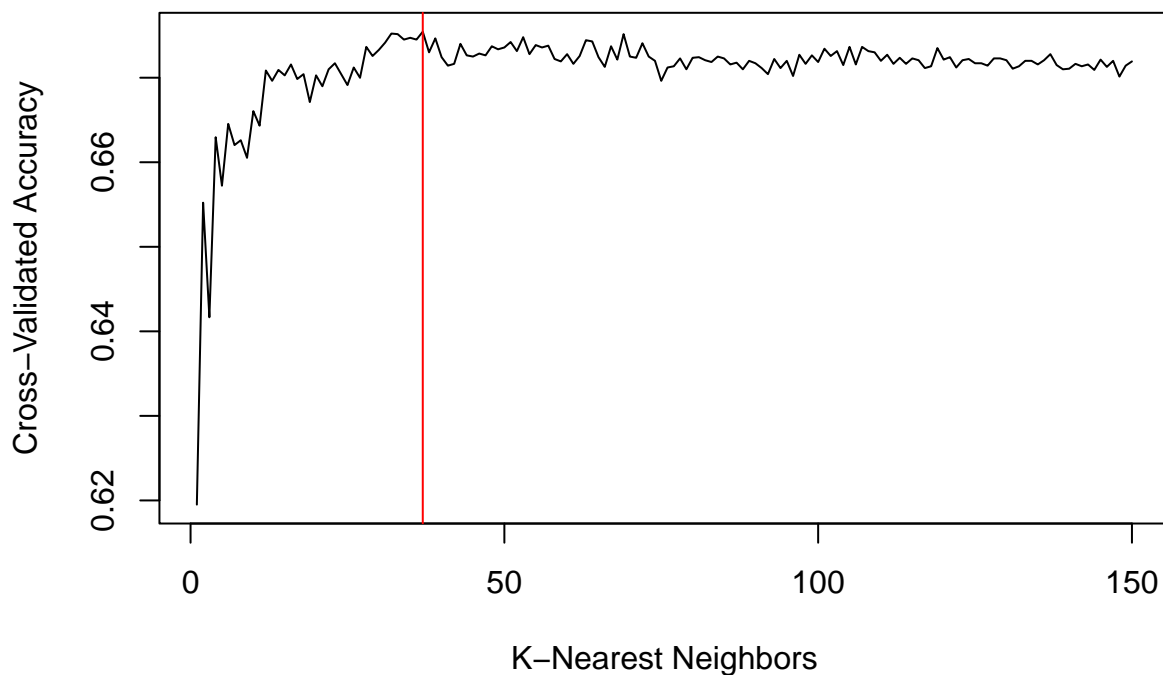
# KNN Train
KS_Scaled = KS_Clean %>%
  mutate_if(is.numeric, scale) %>%
  select(-country, -category, -deadline_weekday, -created_at_weekday,
    -launched_at_weekday, -deadline_month, -deadline_yr,
    -created_at_month, -created_at_yr, -launched_at_yr,
    -launched_at_month, -success)

KS_Train_Scaled = KS_Scaled[Train_Ind,]
KS_Test_Scaled = KS_Scaled[-Train_Ind,]

k = 1:150
KNN_Train = Rfast::knn.cv(nfolds = 10,
  seed = seed,
  y = as.factor(KS_Train$success),
  x = as.matrix(KS_Train_Scaled),
  k = k,
  type = "C",
  pred.ret = TRUE)

optimal_k = which.max(KNN_Train$crit)
plot(k, KNN_Train$crit, type = "l",
  xlab = "K-Nearest Neighbors", ylab = "Cross-Validated Accuracy")
abline(v = optimal_k, col = "red")

```



```
# KNN
KNN = Rfast::knn(xnew = as.matrix(KS_Test_Scaled),
                 y = as.factor(KS_Train$success),
                 x = as.matrix(KS_Train_Scaled),
                 k = optimal_k)-1
```

```
# KNN Performance
```

```
KNN_Table = table(KNN, KS_Test$success)
KNN_Table_Prop = prop.table(KNN_Table)
KNN_Misc = 1-sum(diag(KNN_Table_Prop))
KNN_Misc
```

```
## [1] 0.3164085
```

```
# Results Table
```

```
Results = data.frame(
  Method = c("Logisitic Regression",
             "Lasso Regression",
             "Ridge Regression",
             "Decision Tree",
             "Bagging",
             "Random Forest",
             "Boosting",
             "SVM Linear",
             "SVM Radial",
             "LDA",
             "QDA",
             "KNN"),
  Misclass_rate = c(Logistic_Misc,
                    Lasso_Misc,
                    Ridge_Misc,
                    Pruned_Tree_Misc,
                    Bag_Misc,
                    RF_Misc,
                    Boost_Misc,
                    SVM_Linear_Misc,
                    SVM_Radial_Misc,
                    LDA_Misc,
                    QDA_Misc,
                    KNN_Misc)
)
```

```
Results
```

```
##           Method Misclass_rate
## 1 Logisitic Regression    0.2507172
## 2   Lasso Regression    0.2498566
## 3   Ridge Regression    0.2908778
## 4   Decision Tree    0.2079748
## 5         Bagging    0.1887550
## 6   Random Forest    0.1864601
## 7         Boosting    0.1858864
## 8         SVM Linear    0.2685026
```

## 9	SVM Radial	0.2530120
## 10	LDA	0.3304647
## 11	QDA	0.6044177
## 12	KNN	0.3164085

```
write.csv(Results, "Results.csv")
```